

Pipeline II

Hazards & Branch Detect


Chixiao Chen

Coutesy by

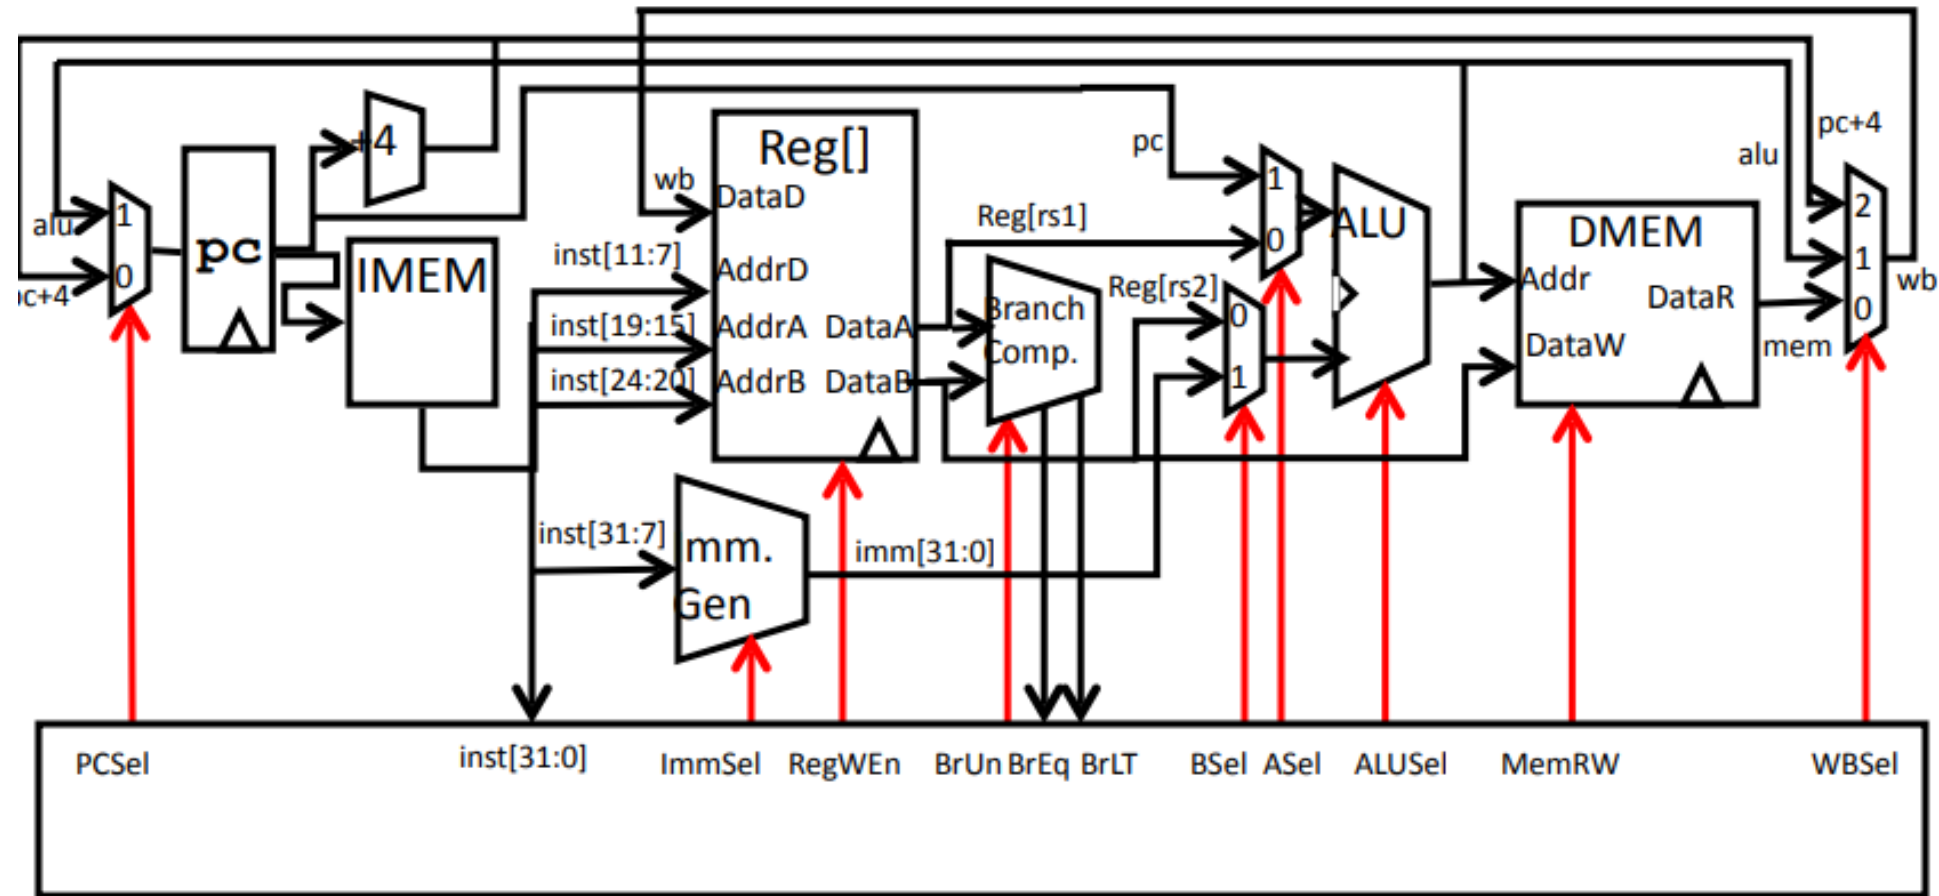
Overview

- Review on the Single Cycle Processor
- Data hazards
 - Why ? Register file timing
 - R-type + R-type
 - Load + R-type
- Control hazard
- AI Perspective

Overall Single Cycle RV32I Datapath


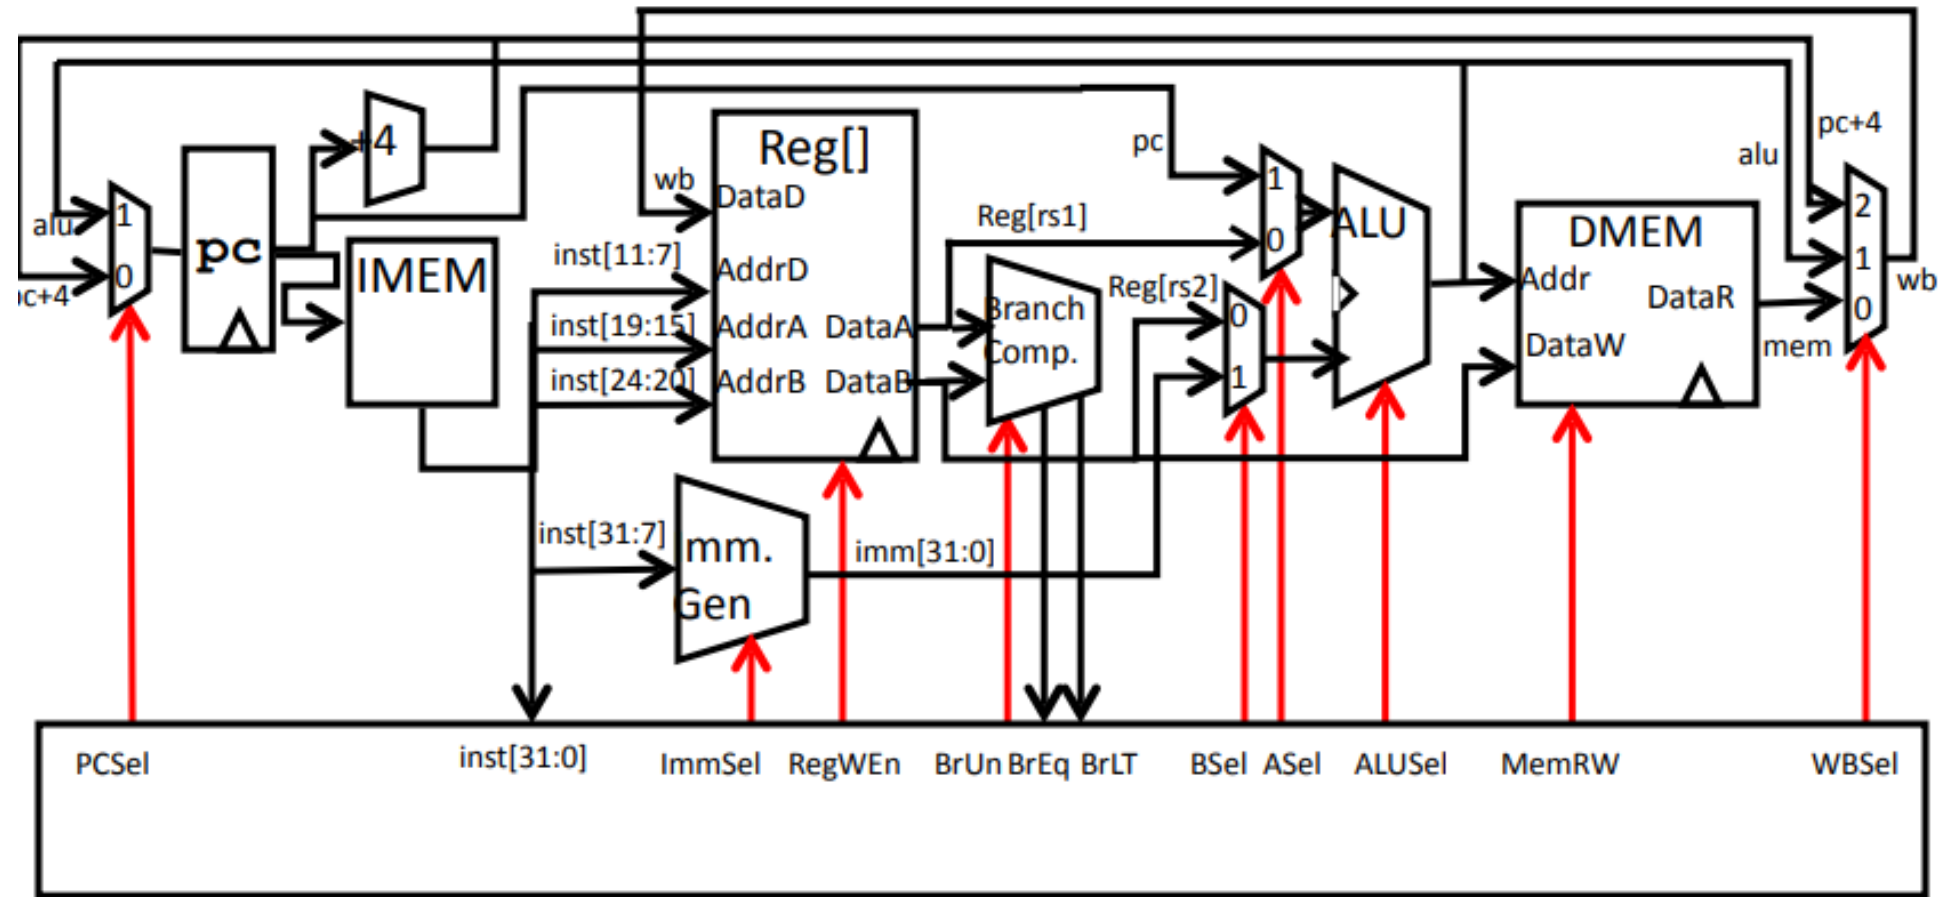


Overall Single Cycle RV32I Datapath

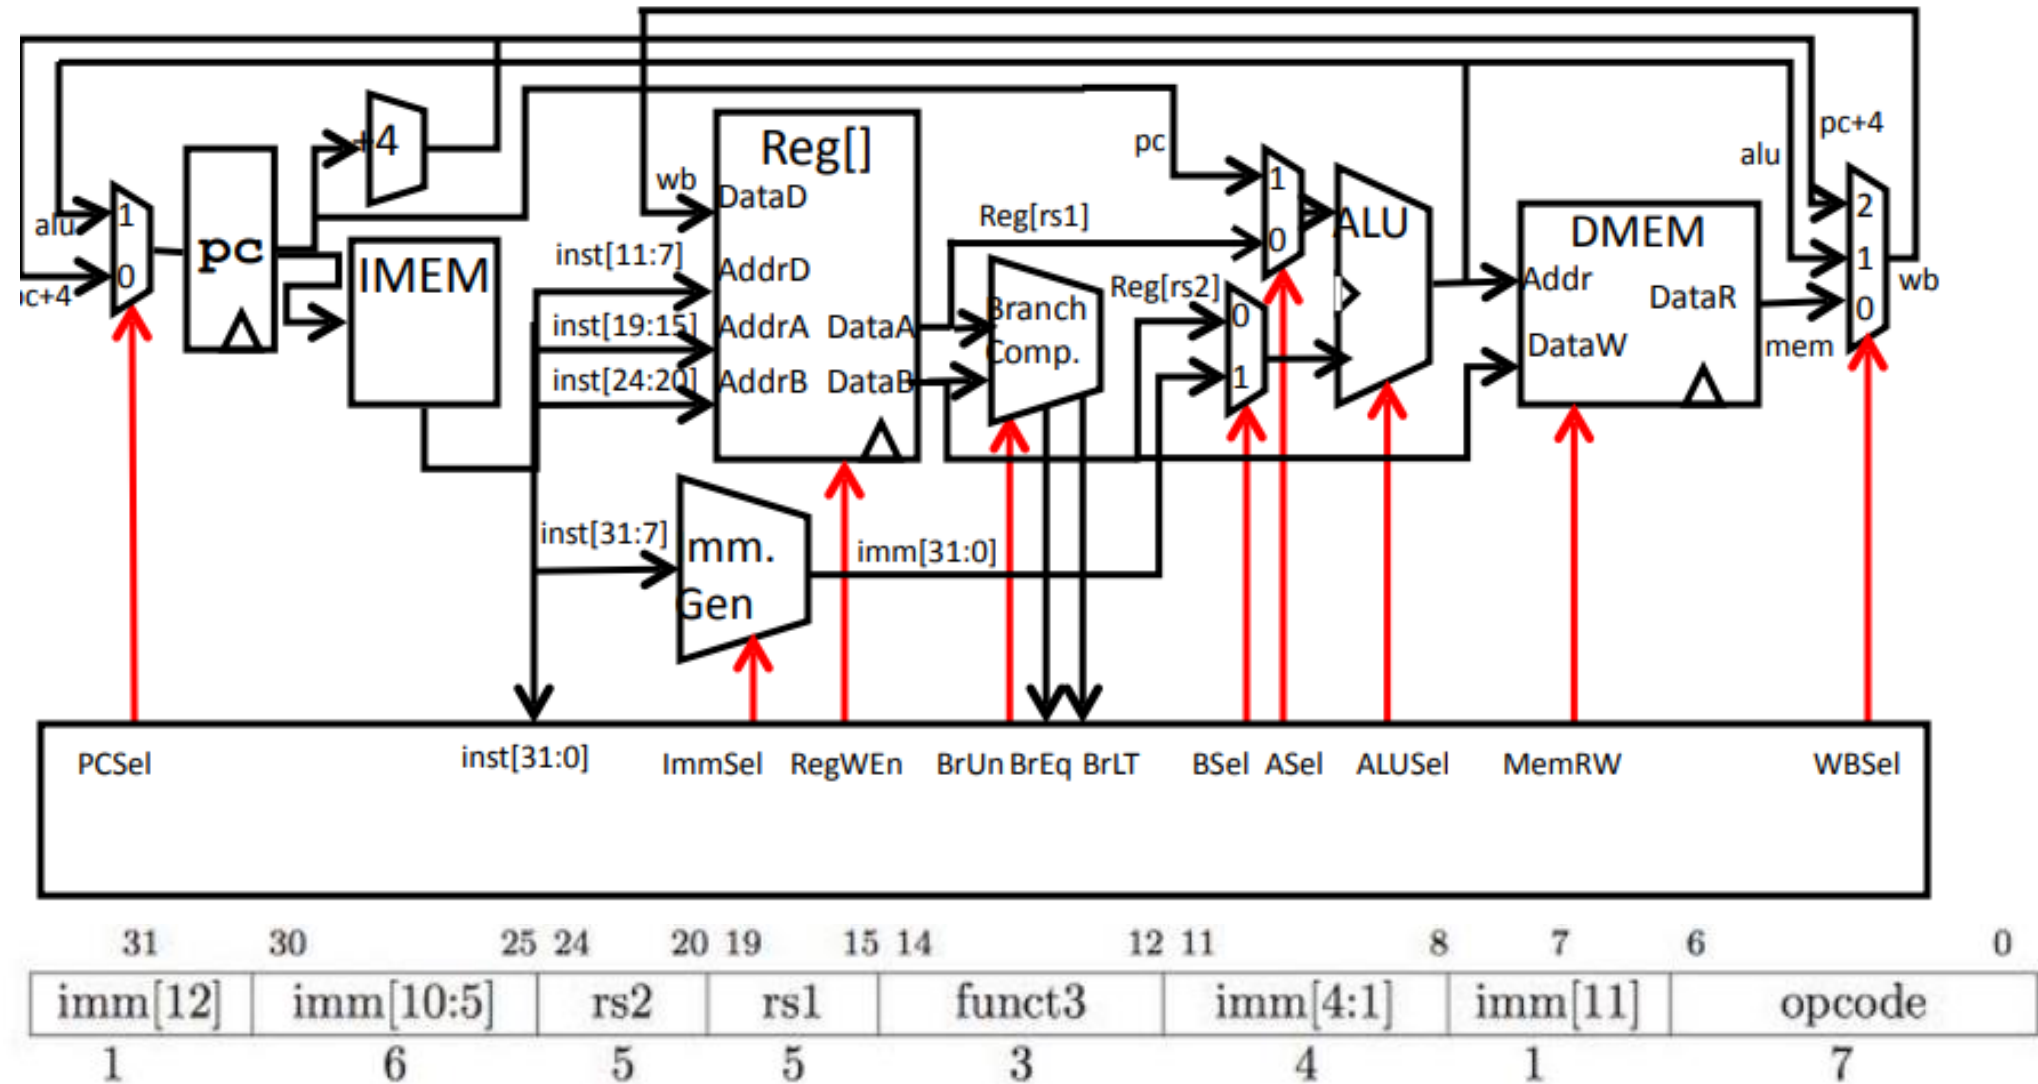


imm[11:0]	rs1	010	rd	0000011	LW
-----------	-----	-----	----	---------	----

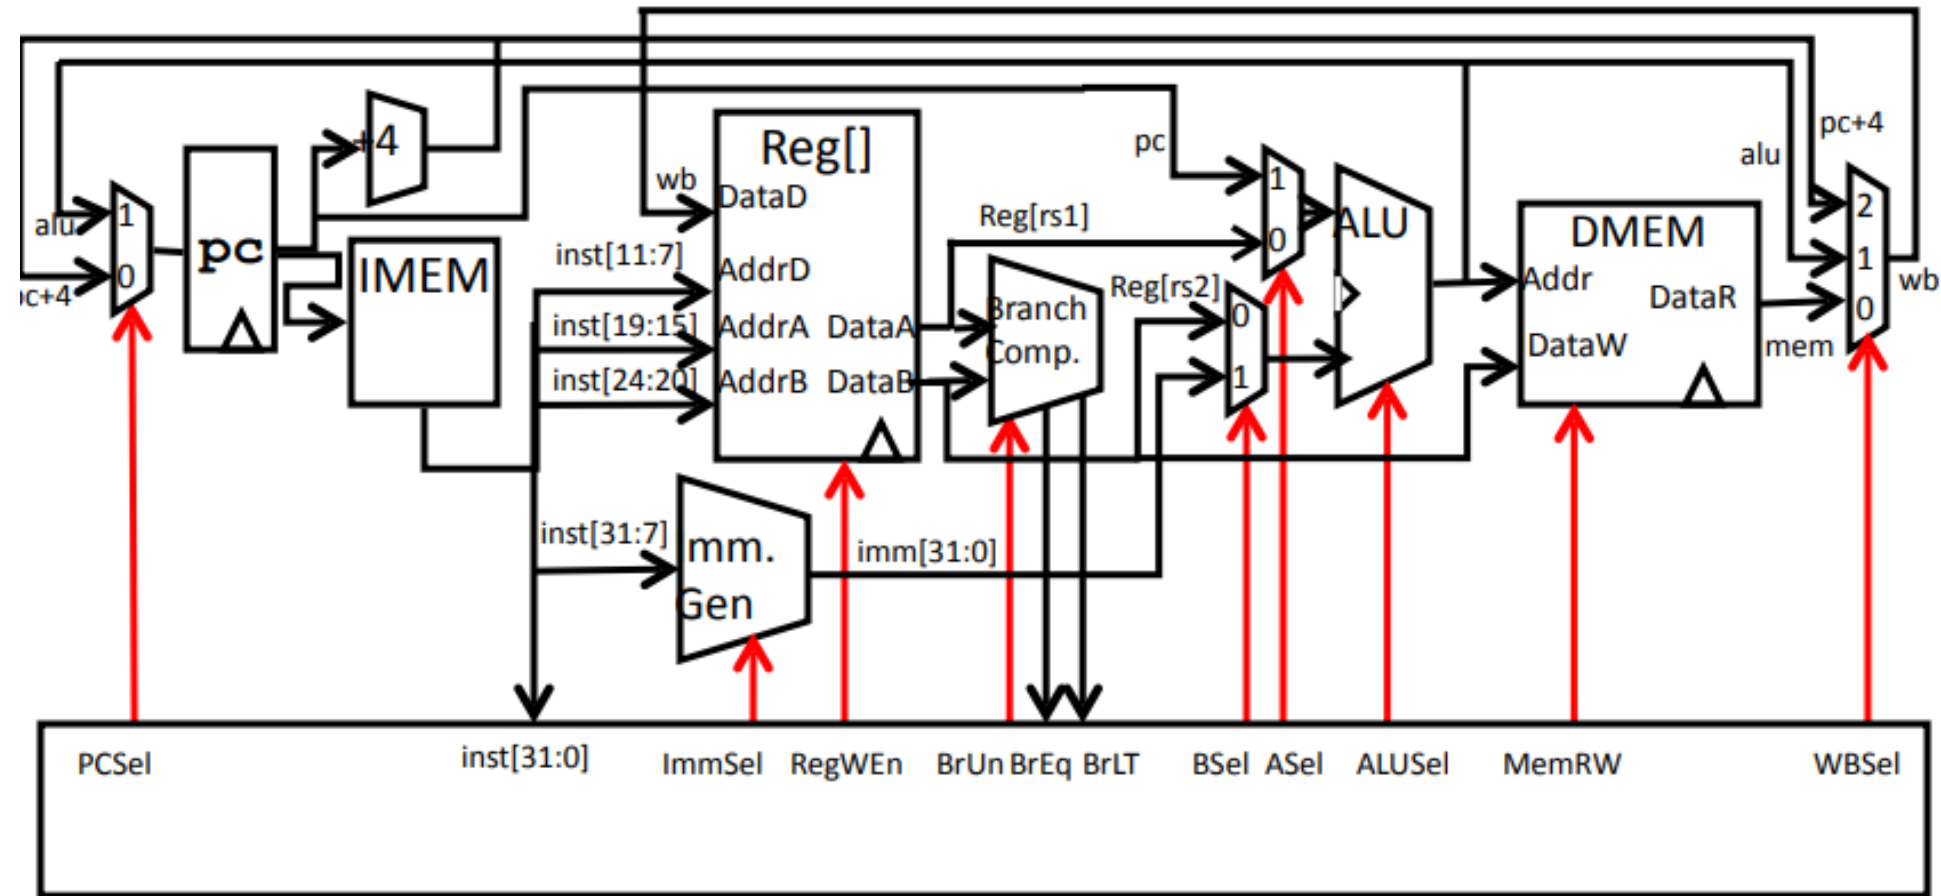
Overall Single Cycle RV32I Datapath



Overall Single Cycle RV32I Datapath



Overall Single Cycle RV32I Datapath



imm[11:0]	rs1	000	rd	1100111	JALR
-----------	-----	-----	----	---------	------


Throughput

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

- Instructions per program depends on source code, compiler technology, and ISA .
- Cycles per instructions (CPI) depends on ISA and μ architecture
- Time per cycle depends upon the μ architecture and base technology


- Why pipeline is good?

Unpipelined machine



3 instructions, 3 cycles, CPI=1

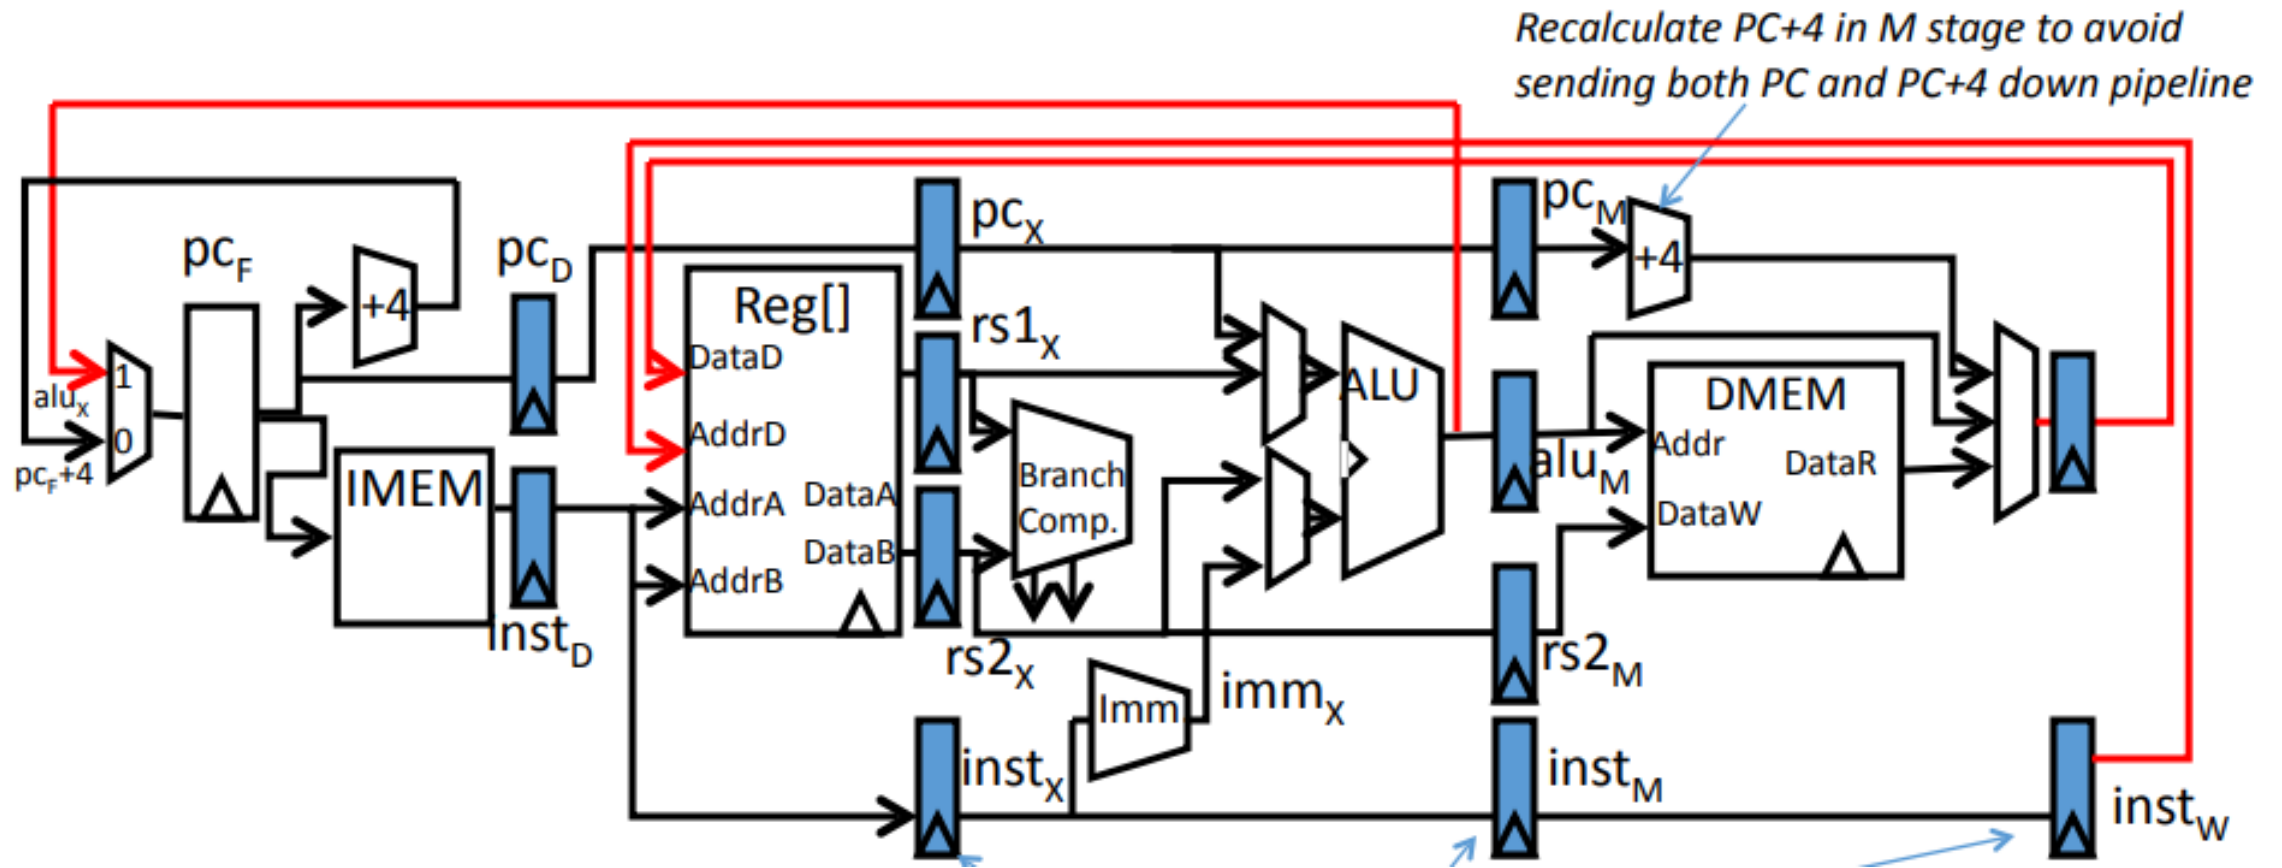
Pipelined machine



3 instructions, 3 cycles, CPI=1

5-stage pipeline CPI=5!!!

Pipelined RV32-I



Hazard

Structural / Data / Control


Hazards Ahead



Hazards

Instructions interact with each other in pipeline

- An instruction in the pipeline may need a resource being used by another instruction in the pipeline → *structural hazard*
- An instruction may depend on something produced by an earlier instruction
 - Dependence may be for a data value
→ *data hazard*
 - Dependence may be for the next instruction's address
→ *control hazard (branches, exceptions)*

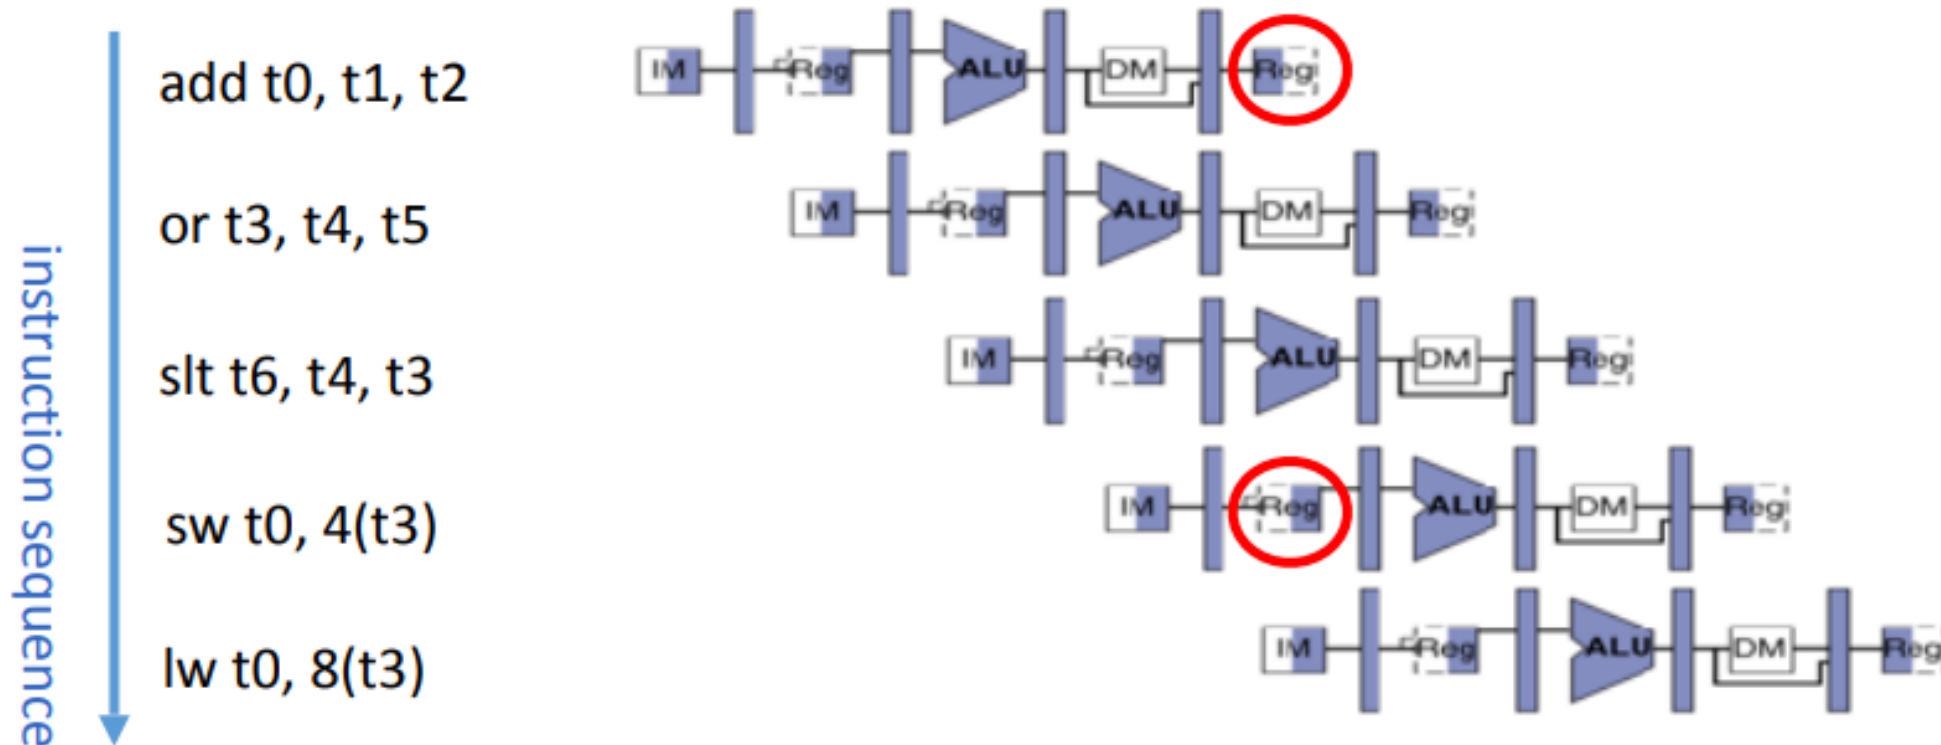


Structural Hazards

- Structural hazard occurs when two instructions need same hardware resource at same time
- A structural hazard can always be avoided by adding more hardware to design
 - E.g., if two instructions both need a port to memory at same time, could avoid hazard by adding second port to memory
- Does Classic RISC 5-stage integer pipeline have structural hazards by design ?

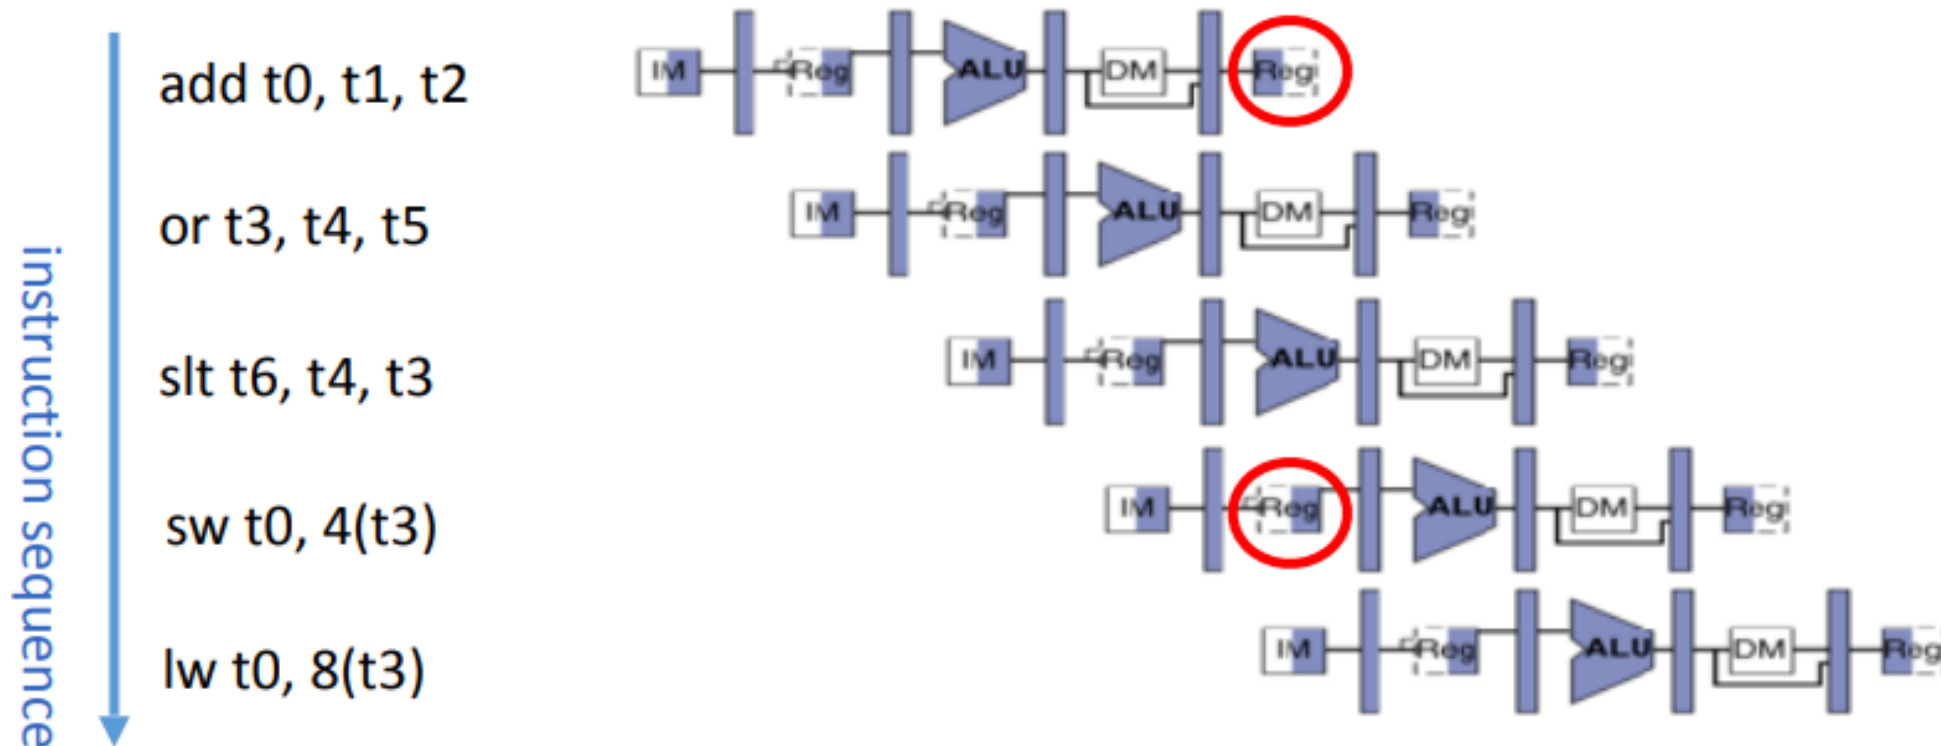
Register Access

- Two instructions using Reg file at same time
- If the data is same, also data hazards



Register Access

- Exploit high speed of register file
 - WB updates value @ posedge / ID reads new value @ negedge



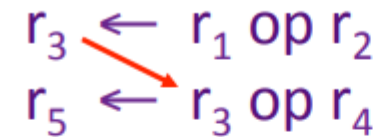
Data Hazards

- Consider executing a sequence of register-register instructions of type:

$$r_k \leftarrow r_i \text{ op } r_j$$

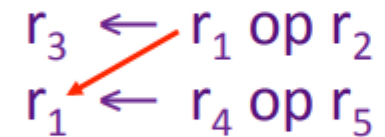
- In 5-stage pipeline, which one shall we care?

Data-dependence

$$\begin{array}{l} r_3 \leftarrow r_1 \text{ op } r_2 \\ r_5 \leftarrow r_3 \text{ op } r_4 \end{array}$$


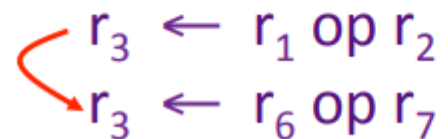
Read-after-Write
(RAW) hazard

Anti-dependence

$$\begin{array}{l} r_3 \leftarrow r_1 \text{ op } r_2 \\ r_1 \leftarrow r_4 \text{ op } r_5 \end{array}$$


Write-after-Read
(WAR) hazard


Output-dependence

$$\begin{array}{l} r_3 \leftarrow r_1 \text{ op } r_2 \\ r_3 \leftarrow r_6 \text{ op } r_7 \end{array}$$


Write-after-Write
(WAW) hazard

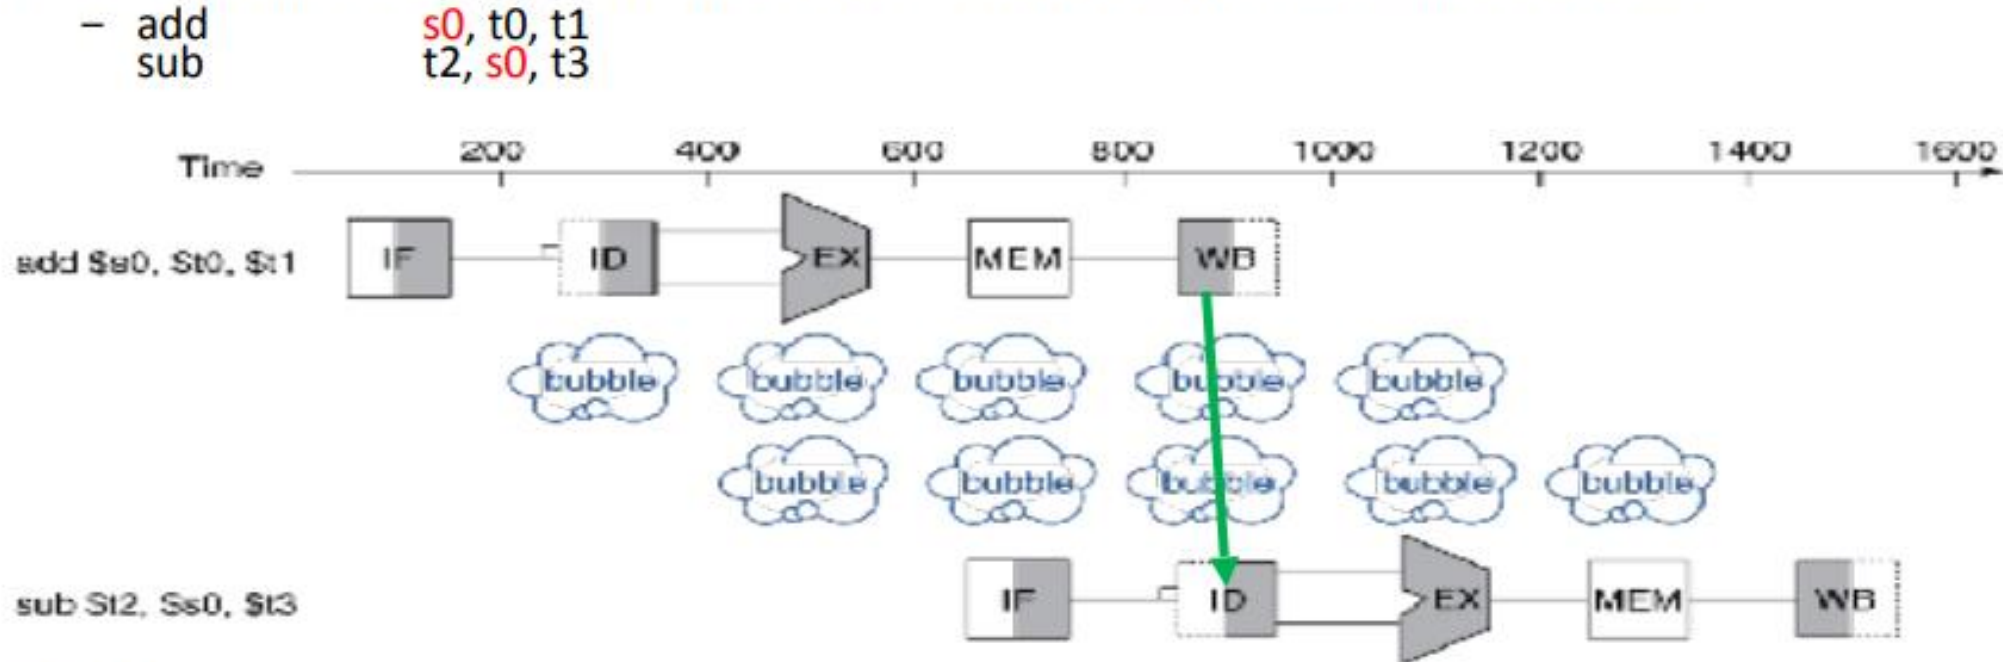
RAW example

s0 holds "5" then add instr changes s0 to "9"



Stall: Simple but not good


- Problem: Instruction depends on result from previous instruction



- Stalls reduce performance
- Compiler could try to arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure


Do we really need to wait ?

- If R-type, the value is calculated after ALU

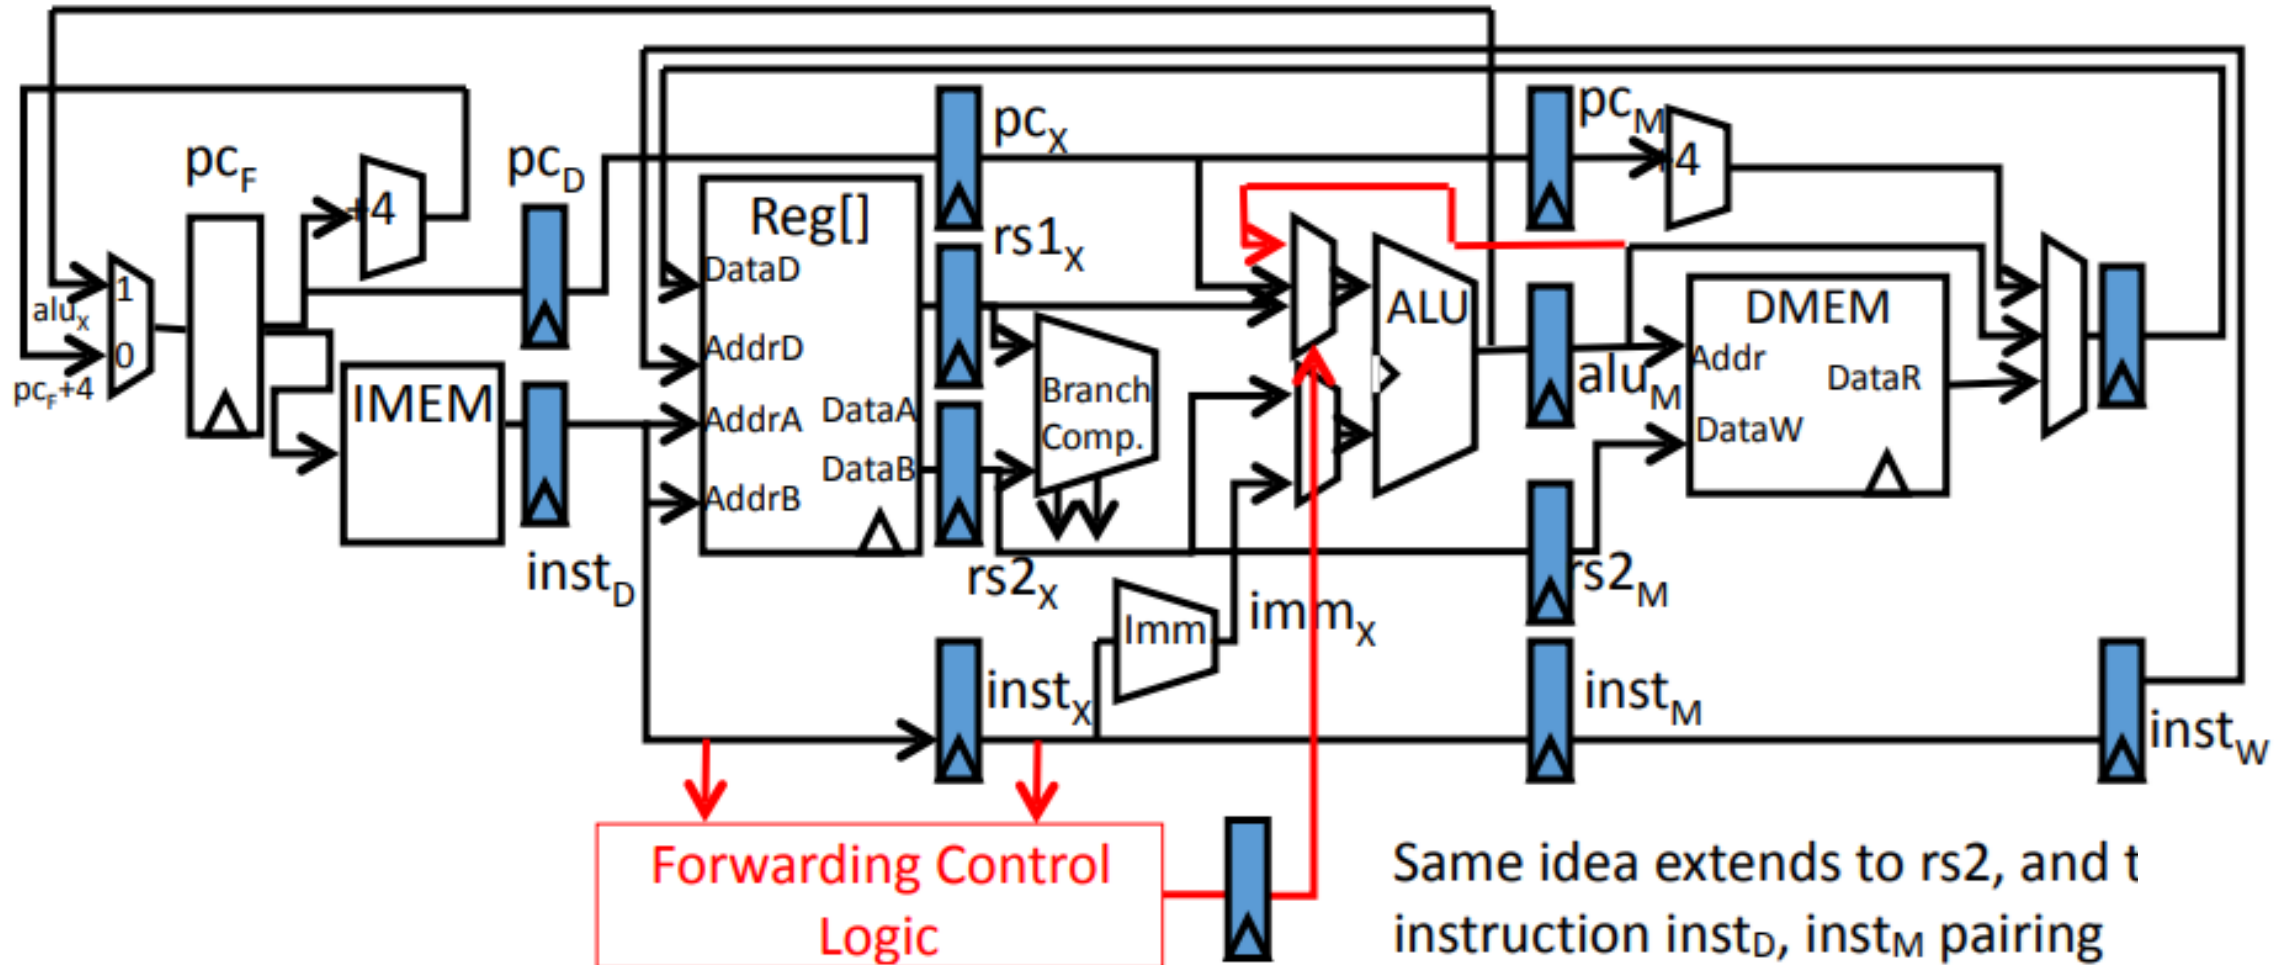


Forwarding / Bypassing

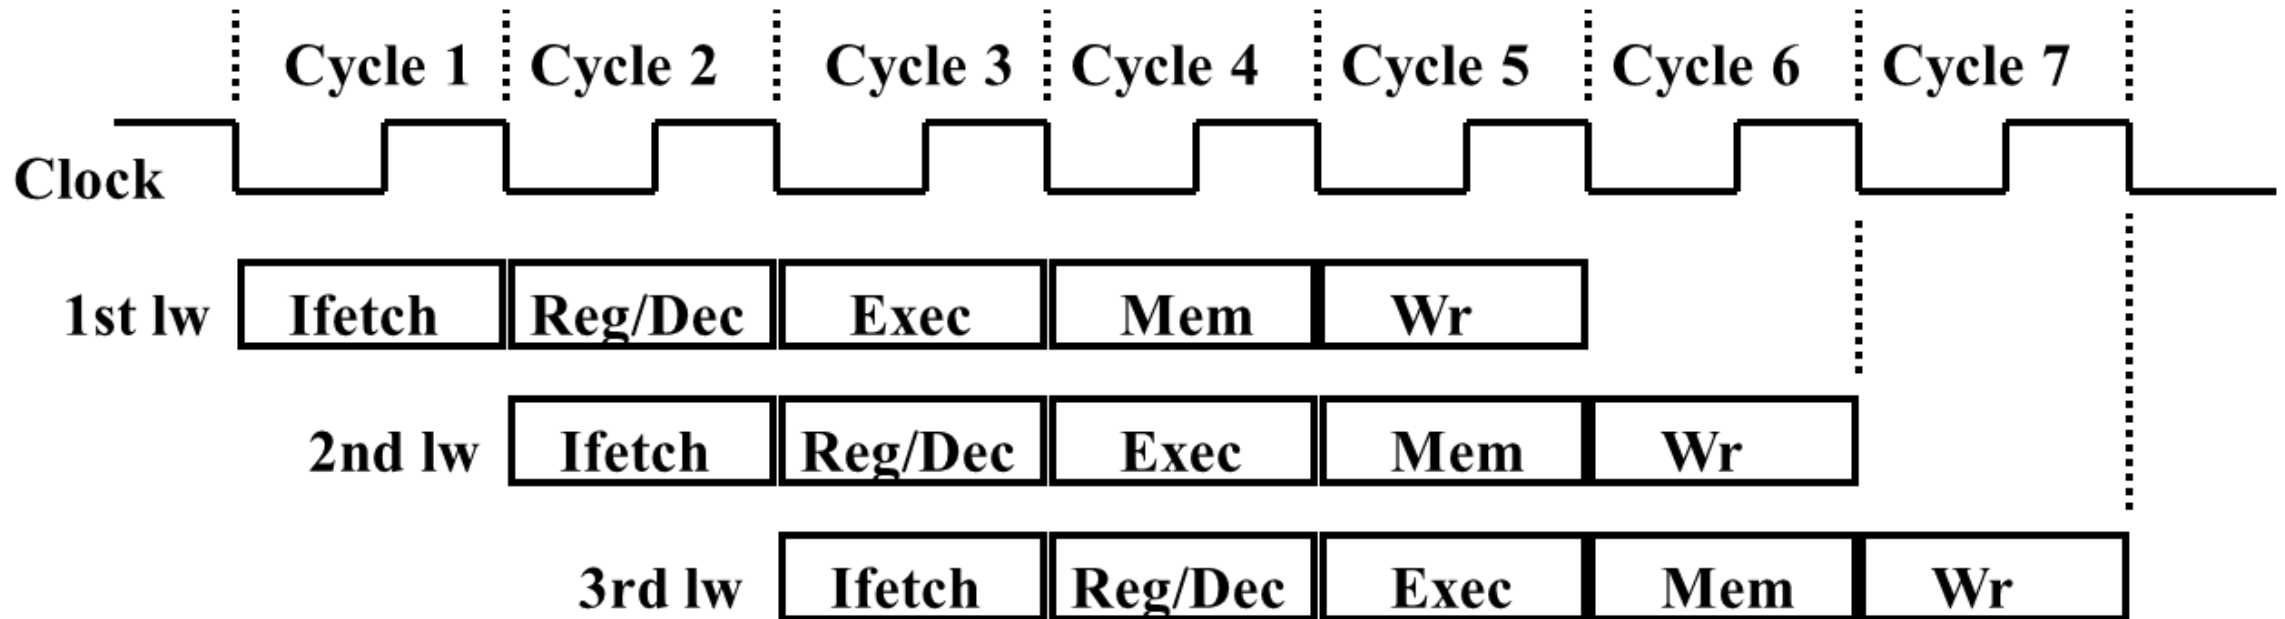
- Use result when it is computed
 - Don't wait for it to be stored in a register



Hardware Modification

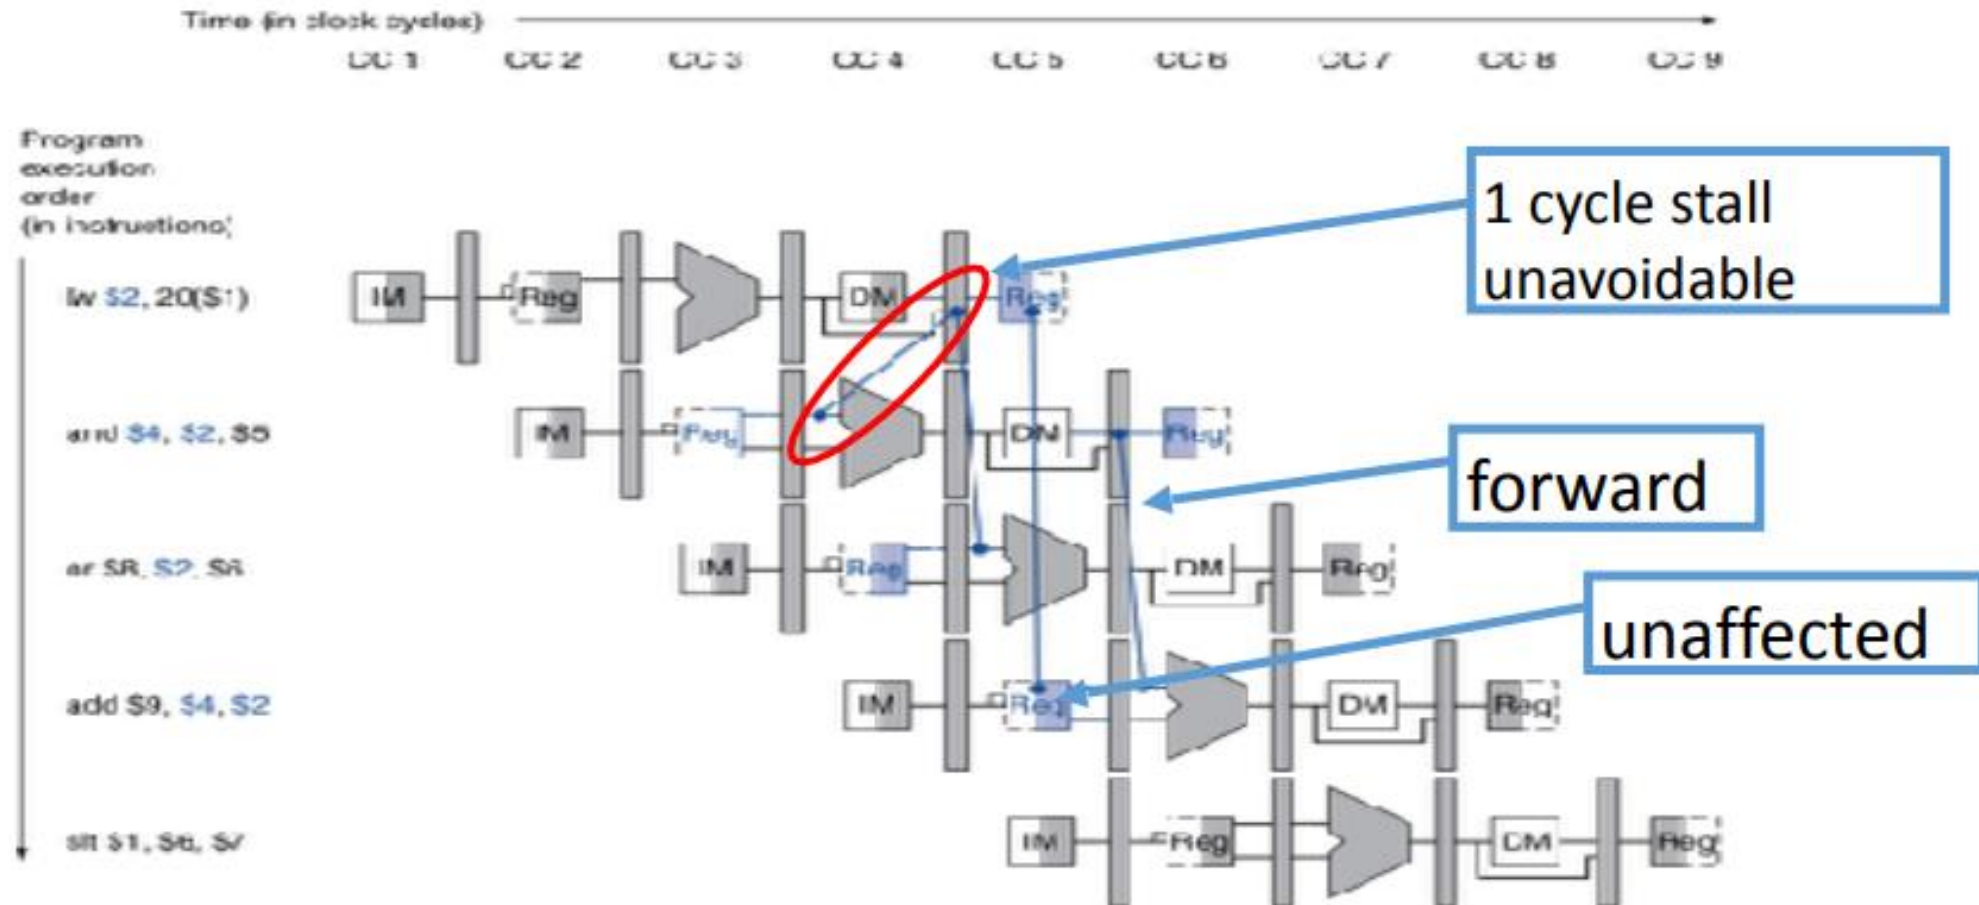


Pipelining load



- 5 functional units in the pipeline datapath are:
 - ALU for the compute the address
 - Data Memory for get data from memery
 - Register File's Write port (busW) for the WB stage

Load data hazard




Stall pipeline

- Slot after a load is called a load delay slot
 - If that instruction uses the result of the load, then the hardware will stall for one cycle
 - Equivalent to inserting an explicit nop in the slot
 - Performance loss!
- Idea:
 - Put unrelated instruction into load delay slot
 - No performance loss!


Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction!
- RISC-V code for $D=A+B$; $E=A+C$;



Control hazard

- When do we need to calculate next PC?
- For Unconditional Jumps
 - Opcode, PC, and offset
- For Jump Register
 - Opcode, Register value, and offset
- For Conditional Branches
 - Opcode, Register (for condition), PC and offset



Control Hazard Example -- Jump

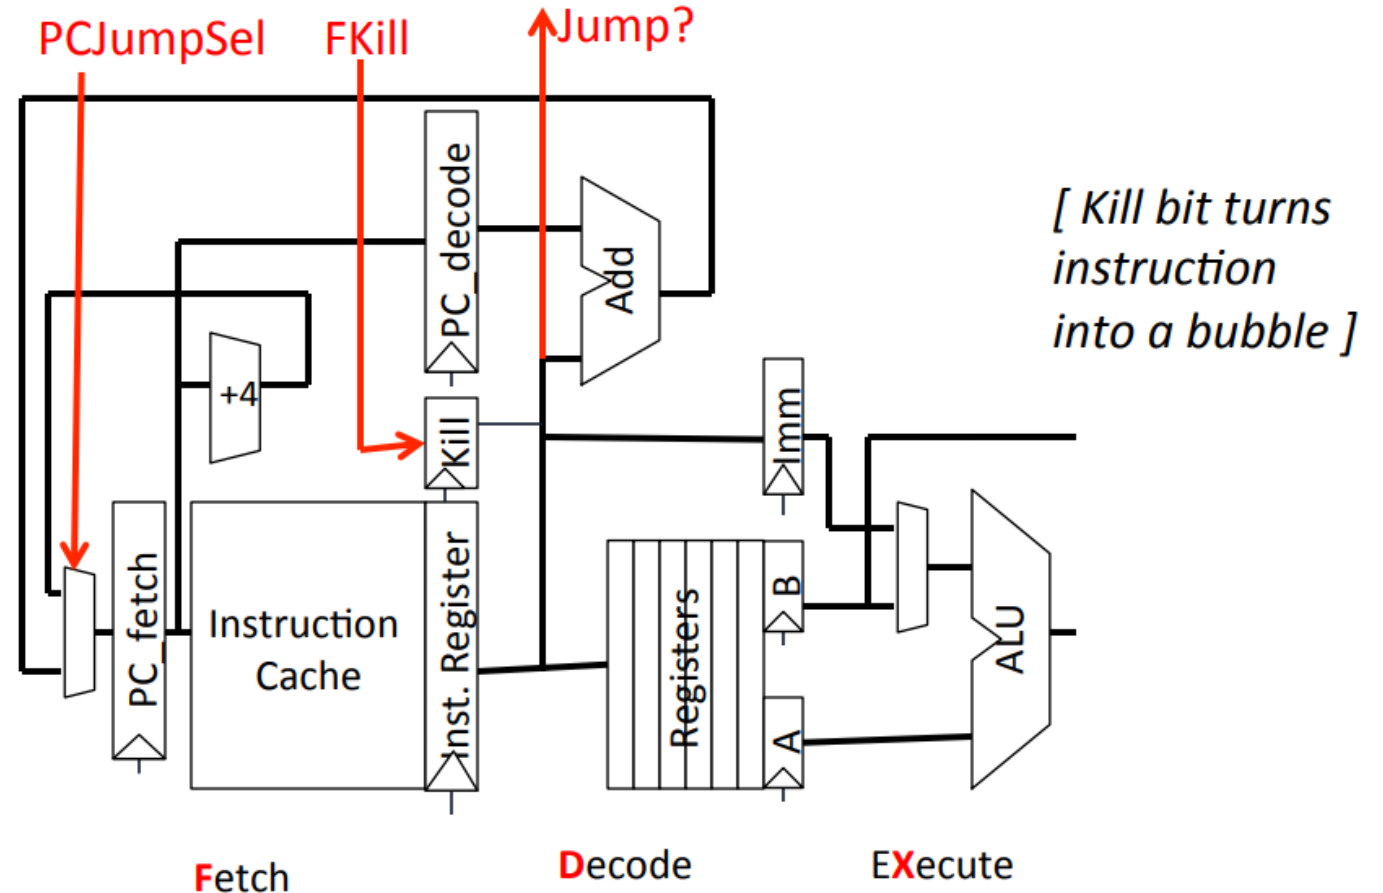
JAL Target

Target: add x1,x2,x3

Hardware modification

Insert kill logic in the control state registers.

Other method?



Control Hazard Example -- Branch

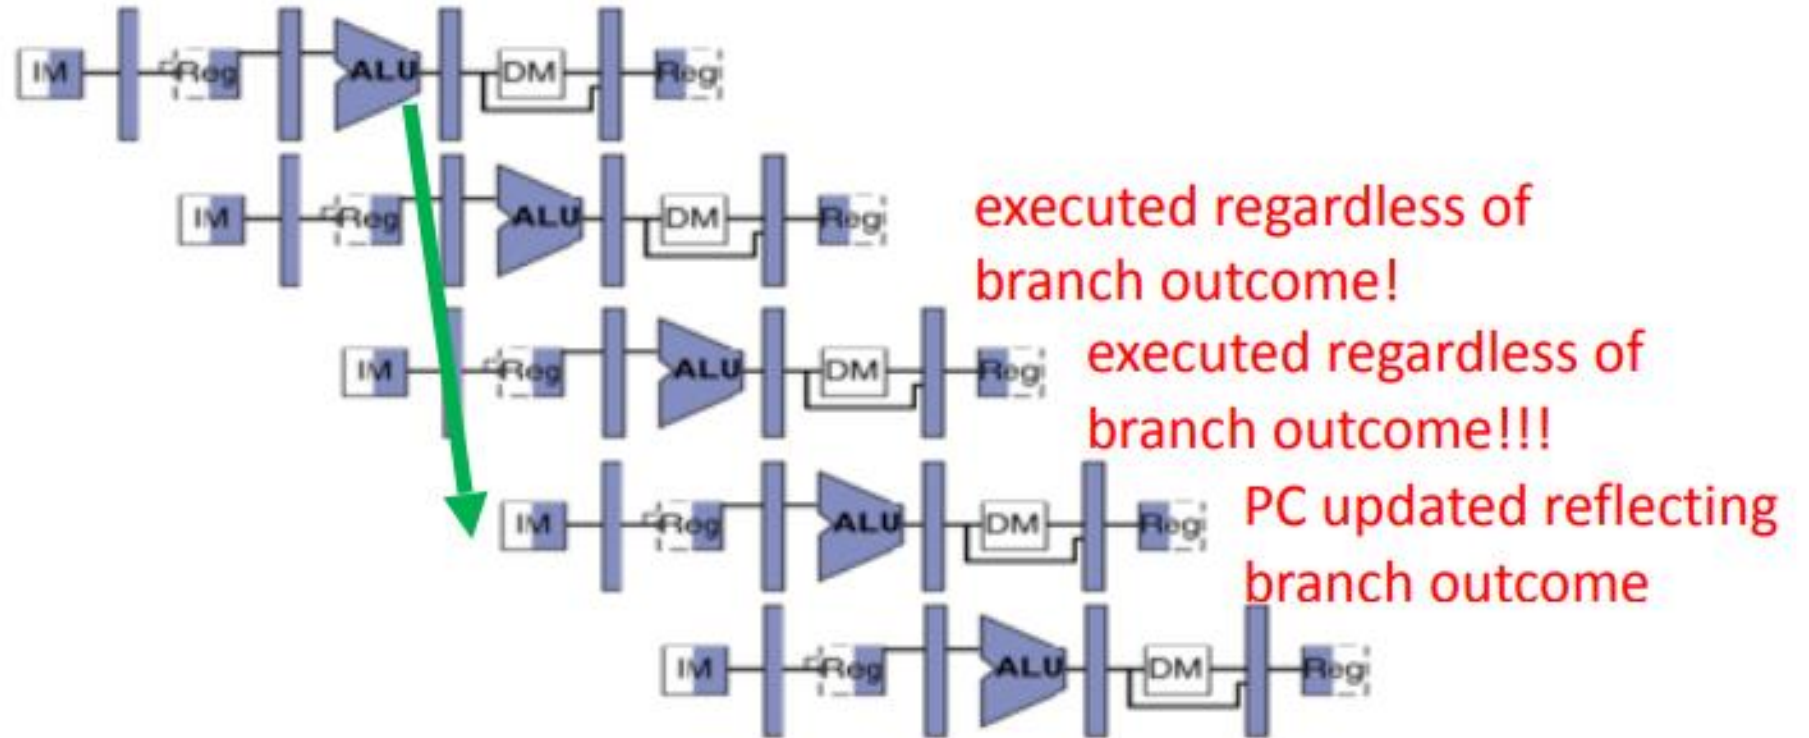
beq t0, t1, label

sub t2, s0, t5

or t6, s0, t3

xor t5, t1, s0

sw s0, 8(t3)




Strategy

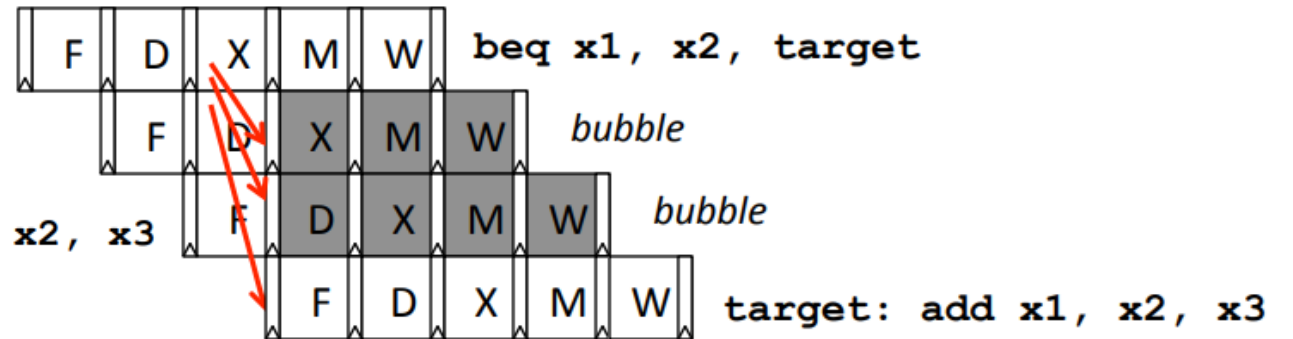
- If branch not taken, then instructions fetched sequentially after branch are correct
- If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs
- How many branches need to be killed in the 5 stage RV32I architecture?

Question: jalr use which style

- JAL

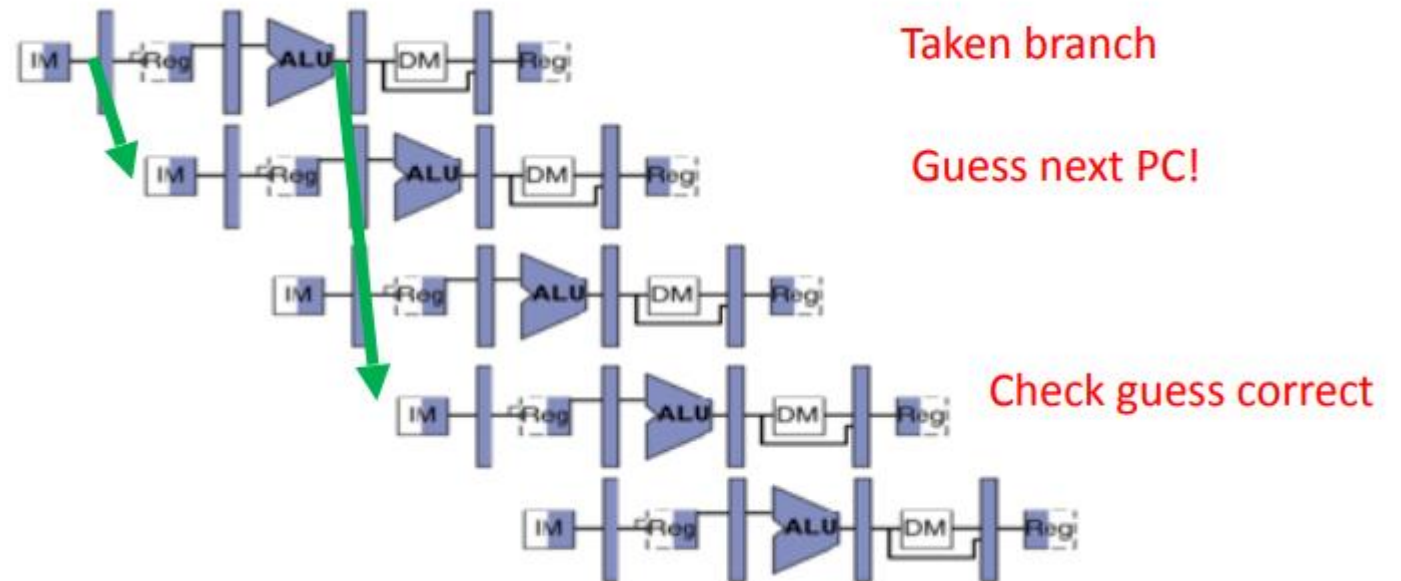


- Branch



Reducing Branch Penalties

- Every taken branch in simple pipeline costs 2 dead cycles
- To improve performance, use “branch prediction” to guess which way branch will go earlier in pipeline
- Only flush pipeline if branch prediction was incorrect




Branch Prediction

- Motivation:
Branch penalties limit performance of deeply pipelined processors
Modern branch predictors have high accuracy (>95%) and can reduce branch penalties significantly
- Required hardware support:
 - Prediction structures
 - Mispredict recovery mechanism

Static Branch Prediction

Overall probability a branch is taken is ~60-70% but:



ISA can attach preferred direction semantics to branches, e.g.,
Motorola MC88110


bne0 (preferred taken) *beq0 (not taken)*

Dynamic Branch Prediction

- Learning based on past behavior
- Temporal correlation
 - The way a branch resolves may be a good predictor of the way it will resolve at the next execution
- Spatial correlation
 - Several branches may resolve in a highly correlated manner (a preferred path of execution)

Temporal Correlation

- One bit history predictor
 - For each branch, remember last way branch went
- Two bit predictor
 - A finite state machine
 - Change the prediction after 2 mistakes



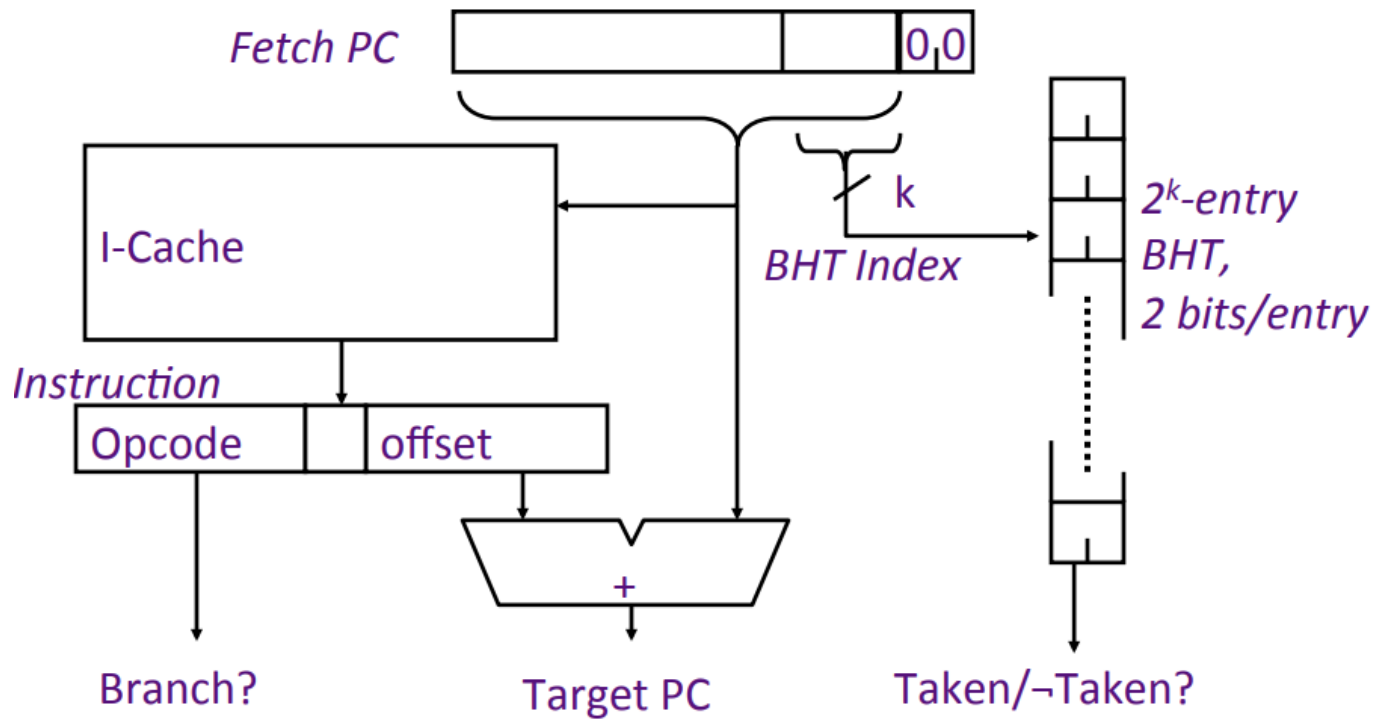
Spatial Correlation

- Other branch style


```
if (aa==2)
    aa=0;
if (bb==2)
    bb=0;
if (aa!=bb) {
    addi x3,x1,-2
    bnez x3,L1      //branch b1  (aa!=2)
    add  x1,x0,x0    //aa=0
L1:  addi x3,x2,-2
    bnez x3,L2      //branch b2  (bb!=2)
    add  x2,x0,x0    //bb=0
L2:  sub  x3,x1,x2    //x3=aa-bb
    beqz x3,L3      //branch b3  (aa==bb)
```

- Prediction can based other branches

Branch History Table

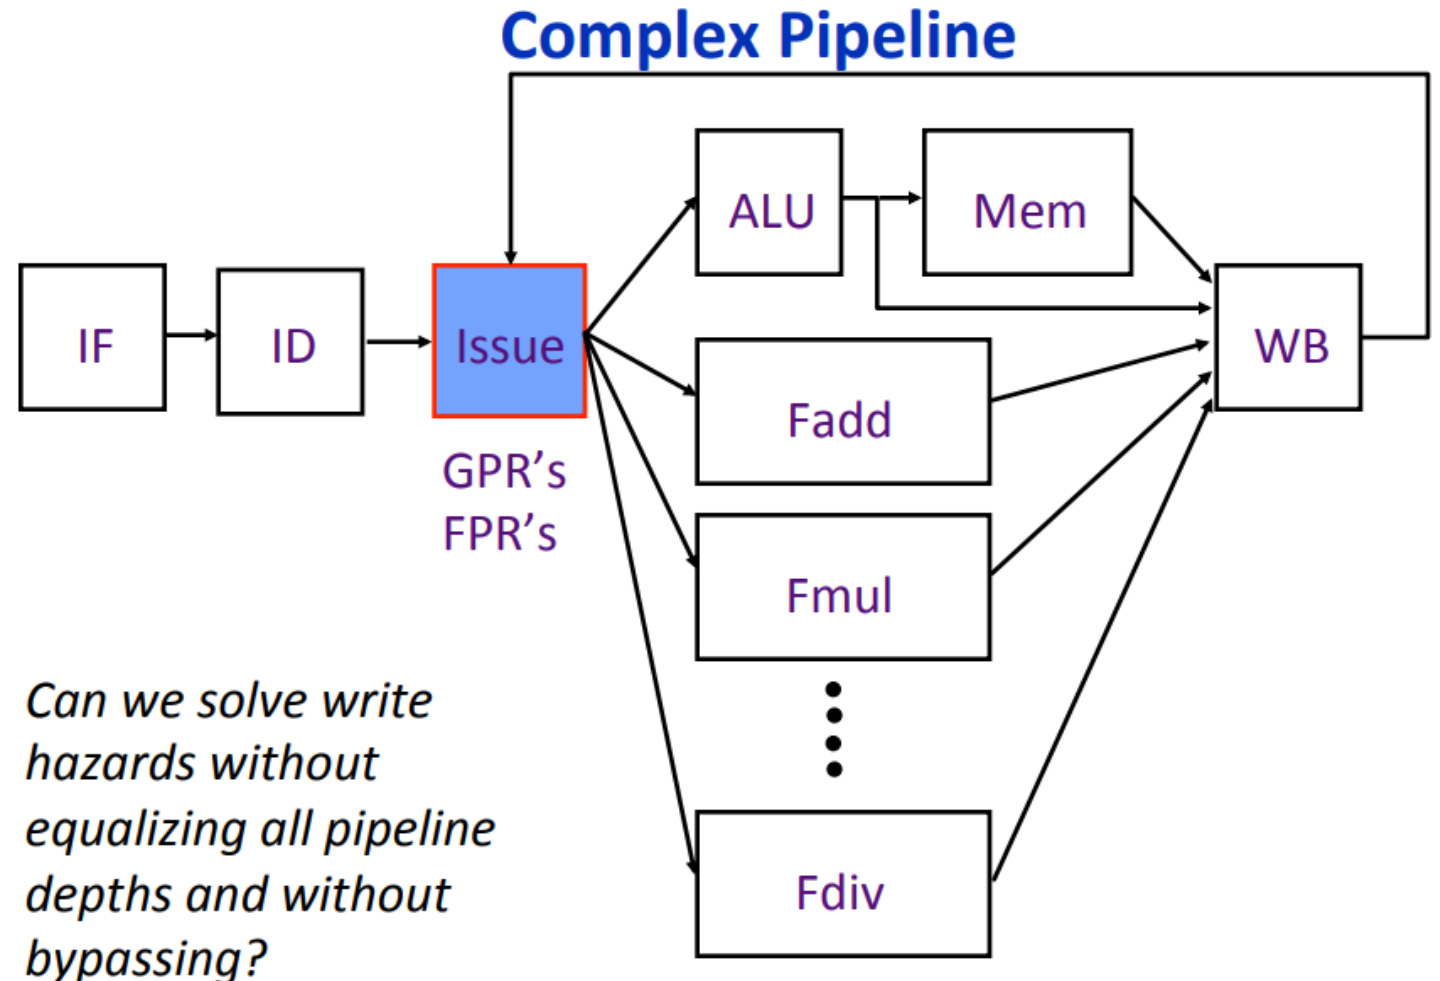


4K-entry BHT, 2 bits/entry, ~80-90% correct predictions



The real world – Complex Pipeline

- Pipeline depths are not equalized in practice.
- Out-of-order scheduling



Loop Unrolling

- There are other methods to avoid branch stalls

```
Loop: fld      f0,0(x1)      //f0=array element
      fadd.d   f4,f0,f2      //add scalar in f2
      fsd      f4,0(x1)      //store result
      addi     x1,x1,-8       //decrement pointer
                                //8 bytes (per DW)
      bne      x1,x2,Loop     //branch x1≠x2
```

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

How many stalls we need?

Without any scheduling, the loop will execute as follows, taking nine cycles:

	<u>Clock cycle issued</u>
Loop: fld f0,0(x1)	1
stall	2
fadd.d f4,f0,f2	3
stall	4
stall	5
fsd f4,0(x1)	6
addi x1,x1,-8	7
bne x1,x2,Loop	8

There are two more after branch instructions

Step 1 – Unrolling the loop

Loop: fld f0,0(x1)	
fadd.d f4,f0,f2	
fsd f4,0(x1)	//drop addi & bne
fld f6,-8(x1)	
fadd.d f8,f6,f2	
fsd f8,-8(x1)	//drop addi & bne
fld f0,-16(x1)	
fadd.d f12,f0,f2	
fsd f12,-16(x1)	//drop addi & bne
fld f14,-24(x1)	
fadd.d f16,f14,f2	
fsd f16,-24(x1)	
addi x1,x1,-32	
bne x1,x2,Loop	

Reschedule the Code



Comparison (normalized to each branch)

- 11 cycles before unrolling
- 6.5 cycles after unrolling
drop addi & bne instruction
- 3.5 cycles after rescheduling

```
Loop: fld    f0,0(x1)
      fld    f6,-8(x1)
      fld    f0,-16(x1)
      fld    f14,-24(x1)
      fadd.d  f4,f0,f2
      fadd.d  f8,f6,f2
      fadd.d  f12,f0,f2
      fadd.d  f16,f14,f2
      fsd     f4,0(x1)
      fsd     f8,-8(x1)
      fsd     f12,16(x1)
      fsd     f16,8(x1)
      addi    x1,x1,-32
      bne     x1,x2,Loop
```


AI Perspective – “Hard” Loop Unrolling

- Flynn Taxonomy
 - SISD: Single Instruction Single Data
 - SIMD: Single Instruction Multiple Data
- More Advanced Topics:
 - Multi-threading
 - SIMT -- GPU



Intel® Architecture currently has SIMD operations of vector length 4, 8, 16

Conclusion

Key words: Pipeline, Hazards, Data Hazard, forwarding, stall, control hazard, branch prediction, loop unrolling, SIMD